

# Rust programming

Module 2: Foundations of Rust

## Unit 3

Advanced Syntax

# Learning objectives

# Composite types

# Types redux

We have previously looked at some of the basic types in the Rust typesystem

- Primitives (integers, floats, booleans, characters)
- Compounds (tuples, arrays)
- Most of the types we looked at were `Copy`
- Borrowing will make more sense when we look at some more ways we can type our data

# Structuring data

Rust has two important ways to structure data

- structs
- enums
- ~~unions~~

# Structs

A struct is similar to a tuple, but this time the combined type gets its own name

```
1 struct ControlPoint(f64, f64, bool);
```

This is an example of a *tuple struct*. You can access the fields in the struct the same way as with tuples:

```
1 fn main() {  
2     let cp = ControlPoint(10.5, 12.3, true);  
3     println!("{}", cp.0); // prints 10.5  
4 }
```

# Structs

Much more common though are structs with named fields

```
1  struct ControlPoint {  
2      x: f64,  
3      y: f64,  
4      enabled: bool,  
5  }
```

- We can add a little more purpose to each field
- No need to keep our indexing up to date when we add or remove a field

```
1  fn main() {  
2      let cp = ControlPoint {  
3          x: 10.5,  
4          y: 12.3,  
5          enabled: true,  
6      };  
7      println!("{}", cp.x); // prints 10.5  
8  }
```

# Enumerations

One of the more powerful kinds of types in Rust are enumerations

```
1  enum IpAddressType {  
2      Ipv4,  
3      Ipv6,  
4  }
```

- An enumeration (listing) of different *variants*
- Each variant is an alternative value of the enum, you pick a single value to create an instance
- Each variant has a discriminant (hidden by default)
  - a numeric value ( ``usize`` by default, can be changed by using ``#[repr(numeric_type)]`` ) used to determine the variant that the enumeration holds
  - one cannot rely on the fact that the discriminant is an ``usize`` , the compiler may always decide to optimize it

```
1  fn main() {  
2      let ip_type = IpAddressType::Ipv4;  
3  }
```



# Enumerations

Enums get more powerful, because each variant can have associated data with it

```
1  enum IPAddress {  
2      Ipv4(u8, u8, u8, u8),           // = 0 (default discriminant)  
3      Ipv6(u16, u16, u16, u16, u16, u16, u16, u16), // = 1 (default discriminant)  
4  }
```

- This way, the associated data and the variant are bound together
- Impossible to create an ipv6 address while only giving a 32 bits integer

```
1  fn main() {  
2      let ipv4_home = IPAddress::Ipv4(127, 0, 0, 1);  
3      let ipv6_home = IPAddress::Ipv6(0, 0, 0, 0, 0, 0, 0, 1);  
4  }
```

- An enum always is as large as the largest variant plus the size of the discriminant

**IPAddress::Ipv4(127, 0, 0, 1)**

<i>isize</i> <small>may be optimized</small>	u8	u8	u8	u8	unused		
---	----	----	----	----	--------	--	--

**IPAddress::Ipv6(0, 0, 0, 0, 0, 0, 0, 1)**

<i>isize</i> <small>may be optimized</small>	u16	u16	u16	u16	u16	u16	u16	u16
---	-----	-----	-----	-----	-----	-----	-----	-----

# Pattern matching

# Extracting data from ``enum``

- We must ensure we interpret ``enum`` data correctly
- Use pattern matching to do so

# Pattern matching

Using the `if let [pattern] = [value]` statement

```
1  fn accept_ipv4(ip: IpAddress) {  
2      if let IpAddress::Ipv4(a, b, _, _) = ip {  
3          println!("Accepted, first octet is {} and second is {}", a, b);  
4      }  
5  }
```

- `a` and `b` introduce local variables within the body of the if that contain the values of those fields
- The underscore ( `_` ) can be used to accept any value

# Match

Pattern matching is very powerful if combined with the match statement

```
1  fn accept_home(ip: IpAddress) {
2      match ip {
3          IpAddress::Ipv4(127, 0, 0, 1) => {
4              println!("You are home!");
5          },
6          IpAddress::Ipv6(0, 0, 0, 0, 0, 0, 0, 1) => {
7              println!("You are in your new home!");
8          },
9          _ => {
10             println!("You are not home");
11         },
12     }
13 }
```

- Every part of the match is called an arm
- A match is exhaustive, meaning all possible values must be handled by one of the match arms
- You can use a catch-all `_` arm to catch any remaining cases if there are any left

# Match as an expression

The match statement can even be used as an expression

```
1  fn get_first_byte(ip: IpAddress) {  
2      let first_byte = match ip {  
3          IpAddress::Ipv4(a, _, _, _) => a,  
4          IpAddress::Ipv6(a, _, _, _, _, _, _, _) => a / 256 as u8,  
5      };  
6      println!("The first byte was: {}", first_byte);  
7  }
```

- The match arms can return a value, but their types have to match
- Note how here we do not need a catch all ( ``_` =>`` ) arm because all cases have already been handled by the two arms

``impl`` blocks

# `impl` blocks

To associate functions to ``structs`` and ``enums``, we use ``impl`` blocks

```
1 fn main() {  
2     let x = Some(42);  
3     let unwrapped = x.unwrap();  
4     println!("{}", unwrapped);  
5 }
```

- The syntax ``x.y()`` looks similar to how we accessed a field in a struct
- We can define functions on our types using impl blocks
- Impl blocks can be defined on any type, not just structs (with some limitations)



## `impl` blocks

```
1  enum IpAddress {
2      Ipv4(u8, u8, u8, u8),
3      Ipv6(u16, u16, u16, u16, u16, u16, u16, u16),
4  }
5
6  impl IpAddress {
7      fn as_u32(&self) -> Option<u32> {
8          match self {
9              IpAddress::Ipv4(a, b, c, d) => a << 24 + b << 16 + c << 8 + d
10             _ => None, _
11         }
12     }
13 }
14
15 fn main() {
16     let addr = IpAddress::Ipv4(127, 0, 0, 1);
17     println!("{:?}", addr.as_u32());
18 }
```

# `self` and `Self`

- The ``self`` parameter defines how the method can be used.
- The ``Self`` type is a shorthand for the type on which the current implementation is specified.

```
1  struct Foo(i32);
2
3  impl Foo {
4      fn consume(self) -> Self {           // Takes `Foo` by value, returns `Foo`
5          Self(self.0 + 1)
6      }
7
8      fn borrow(&self) -> &i32 {           // Takes immutable reference of `Foo`
9          &self.0
10     }
11
12     fn borrow_mut(&mut self) -> &mut i32 { // Takes mutable reference of `Foo`
13         &mut self.0
14     }
15
16     fn new() -> Self {                   // Associated function, returns `Foo`
17         Self(0)
18     }
19 }
```

# `impl` blocks, the `self` parameter

The `self` parameter is called the *receiver*.

- The `self`` parameter is always the first and it always has the type on which it was defined
- We never specify the type of the `self`` parameter
- We can optionally prepend `&`` or `&mut`` to `self`` to indicate that we take a value by reference
- Absence of a `self`` parameter means that the function is an associated function instead

```
1 fn main () {  
2     let mut f = Foo::new();  
3     println!("{}", f.borrow());  
4     *f.borrow_mut() = 10;  
5     let g = f.consume();  
6     println!("{}", g.borrow());  
7 }
```

# Optionals and Error handling

# Generics

Structs become even more powerful if we introduce a little of generics

```
1 struct PointFloat(f64, f64);  
2 struct PointInt(i64, i64);
```

We are repeating ourselves here, what if we could write a data structure for both of these cases?

```
1 struct Point<T>(T, T);  
2  
3 fn main() {  
4     let float_point: Point<f64> = Point(10.0, 10.0);  
5     let int_point: Point<i64> = Point(10, 10);  
6 }
```

Generics are much more powerful, but this is all we need for now

# Option

A quick look into the basic enums available in the standard library

- Rust does not have null, but you can still define variables that optionally do not have a value
- For this you can use the `Option<T>` enum

```
1  enum Option<T> {  
2      Some(T),  
3      None,  
4  }  
5  
6  fn main() {  
7      let some_int = Option::Some(42);  
8      let no_string: Option<String> = Option::None;  
9  }
```

# Option

A quick look into the basic enums available in the standard library

- Rust does not have null, but you can still define variables that optionally do not have a value
- For this you can use the `Option<T>` enum

```
1  enum Option<T> {  
2      Some(T),  
3      None,  
4  }  
5  
6  fn main() {  
7      let some_int = Some(42);  
8      let no_string: Option<String> = None;  
9  }
```

# Error handling

What would we do when there is an error?

```
1  fn divide(x: i64, y: i64) -> i64 {  
2      if y == 0 {  
3          // what to do now?  
4      } else {  
5          x / y  
6      }  
7  }
```



# Error handling

What would we do when there is an error?

```
1  fn divide(x: i64, y: i64) -> i64 {  
2      if y == 0 {  
3          panic!("Cannot divide by zero");  
4      } else {  
5          x / y  
6      }  
7  }
```

- A panic in Rust is the most basic way to handle errors
- A panic error is an all or nothing kind of error
- A panic will immediately stop running the current thread/program and instead immediately work to shut it down, using one of two methods:
  - Unwinding: going up through the stack and making sure that each value is cleaned up
  - Aborting: ignore everything and immediately exit the thread/program
- Only use panic in small programs if normal error handling would also exit the program
- Avoid using panic in library code or other reusable components

# Error handling

What would we do when there is an error? We could try and use the option enum instead of panicking

```
1  fn divide(x: i64, y: i64) -> Option<i64> {  
2      if y == 0 {  
3          None  
4      } else {  
5          Some(x / y)  
6      }  
7  }
```

# Result

Another really powerful enum is the result, which is even more useful if we think about error handling

```
1  enum Result<T, E> {
2      Ok(T),
3      Err(E),
4  }
5
6  enum DivideError {
7      DivisionByZero,
8      CannotDivideOne,
9  }
10
11 fn divide(x: i64, y: i64) -> Result<i64, DivideError> {
12     if x == 1 {
13         Err(DivideError::CannotDivideOne)
14     } else if y == 0 {
15         Err(DivideError::DivisionByZero)
16     } else {
17         Ok(x / y)
18     }
19 }
```

# Handling results

Now that we have a function that returns a result we have to think about how we handle that error at the call-site

```
1 fn div_zero_fails() {  
2     match divide(10, 0) {  
3         Ok(div) => println!("{}", div),  
4         Err(e) => panic!("Could not divide by zero"),  
5     }  
6 }
```

- We made the signature of the `divide` function explicit in how it can fail
- The user of the function can now decide what to do, even if it is panicking
- Note: just as with `Option` we never have to use `Result::Ok` and `Result::Err` because they have been made available globally

# Handling results

Especially when writing initial prototyping code you will often find yourself wanting to write error handling code later, Rust has a useful utility function to help you for both `Option` and `Result` :

```
1 fn div_zero_fails() {  
2     let div = divide(10, 0).unwrap();  
3     println!("{}", div);  
4 }
```

- Unwrap checks if the Result/Option is `Ok(x)` or `Some(x)` respectively and then return that `x`, otherwise it will panic your program with an error message
- Having unwraps all over the place is generally considered a bad practice
- Sometimes you can ensure that an error won't occur, in such cases `unwrap` can be a good solution

# Handling results

Especially when writing initial prototyping code you will often find yourself wanting to write error handling code later, Rust has a useful utility function to help you for both `Option` and `Result` :

```
1 fn div_zero_fails() {  
2     let div = divide(10, 0).unwrap_or(-1);  
3     println!("{}", div);  
4 }
```

Besides `unwrap`, there are some other useful utility functions

- `unwrap_or(val)` : If there is an error, use the value given to `unwrap_or` instead
- `unwrap_or_default()` : Use the default value for that type if there is an error
- `expect(msg)` : Same as `unwrap`, but instead pass a custom error message
- `unwrap_or_else(fn)` : Same as `unwrap_or`, but instead call a function that generates a value in case of an error

# Result and the `?`` operator

Results are so common that there is a special operator associated with them, the `?`` operator

```
1 fn can_fail() -> Result<i64, DivideError> {
2     let intermediate_result = match divide(10, 0) {
3         Ok(ir) => ir,
4         Err(e) => return Err(e),
5     };
6
7     match divide(intermediate_result, 0) {
8         Ok(sec) => Ok(sec * 2),
9         Err(e) => Err(e),
10    }
11 }
```

Look how this function changes if we use the `?`` operator

```
1 fn can_fail() -> Result<i64, DivideError> {
2     let intermediate_result = divide(10, 0)?;
3     Ok(divide(intermediate_result, 0)? * 2)
4 }
```

# Result and the `?`` operator

```
1 fn can_fail() -> Result<i64, DivideError> {  
2     let intermediate_result = divide(10, 0)?;  
3     Ok(divide(intermediate_result, 0)? * 2)  
4 }
```

- The `?`` operator does an implicit match, if there is an error, that error is then immediately returned and the function returns early
- If the result is `Ok()`` then the value is extracted and we can continue right away



``Vec``

## `Vec`: storing more of the same

The vector is an array that can grow

- Compare this to the array we previously saw, which has a fixed size

```
1  fn main() {  
2      let arr = [1, 2];  
3      println!("{:?}", arr);  
4  
5      let mut nums = Vec::new();  
6      nums.push(1);  
7      nums.push(2);  
8      println!("{:?}", nums);  
9  }
```

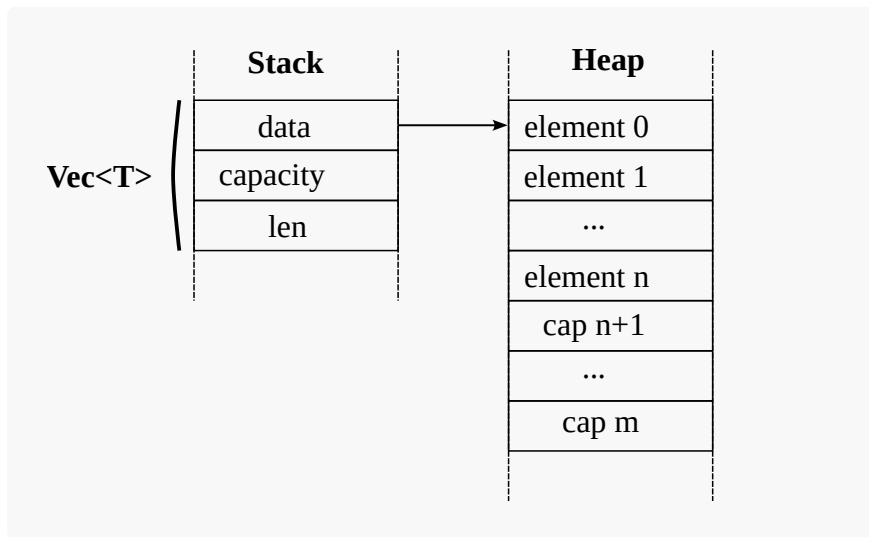
# `Vec`

`Vec` is such a common type that there is an easy way to initialize it with values that looks similar to arrays

```
1 fn main() {  
2     let mut nums = vec![1, 2];  
3     nums.push(3);  
4     println!("{:?}", nums);  
5 }
```

# `Vec`: memory layout

How can a vector grow? Things on the stack need to be of a fixed size



Slices

# Vectors and arrays

What if we wanted to write a sum function, we could define one for arrays of a specific size:

```
1  fn sum(data: &[i64; 10]) -> i64 {  
2      let mut total = 0;  
3      for val in data {  
4          total += val;  
5      }  
6      total  
7  }
```

# Vectors and arrays

Or one for just vectors:

```
1  fn sum(data: &Vec<i64>) -> i64 {  
2      let mut total = 0;  
3      for val in data {  
4          total += val;  
5      }  
6      total  
7  }
```

# Slices

What if we want something to work on arrays of any size? Or what if we want to support summing up only parts of a vector?

- A slice is a dynamically sized view into a contiguous sequence
- Contiguous: elements are laid out in memory such that they are evenly spaced
- Dynamically sized: the size of the slice is not stored in the type, but is determined at runtime
- View: a slice is never an owned data structure
- Slices are typed as `[T]`, where `T` is the type of the elements in the slice



# Slices

```
1 fn sum(data: [i64]) -> i64 {
2     let mut total = 0;
3     for val in data {
4         total += val;
5     }
6     total
7 }
8
9 fn main() {
10     let data = vec![10, 11, 12, 13, 14];
11     println!("{}", sum(data));
12 }
```

```
1     Compiling playground v0.0.1 (/playground)
2 error[E0277]: the size for values of type `[i64]` cannot be known at compilation time
3 --> src/main.rs:1:8
4   |
5 1 | fn sum(data: [i64]) -> i64 {
6   |             ^^^^^ doesn't have a size known at compile-time
7   |
8   = help: the trait `Sized` is not implemented for `[i64]`
9 help: function arguments must have a statically known size, borrowed types always have a known size
```


# Slices

```
1  fn sum(data: &[i64]) -> i64 {
2      let mut total = 0;
3      for val in data {
4          total += val;
5      }
6      total
7  }
8
9  fn main() {
10     let data = vec![10, 11, 12, 13, 14];
11     println!("{}", sum(&data));
12 }
```

```
1  Compiling playground v0.0.1 (/playground)
2  Finished dev [unoptimized + debuginfo] target(s) in 0.89s
3  Running `target/debug/playground`
4  60
```

# Slices

- `[T]` is an incomplete type: we need to know how many `T` there are
- Types that have a known compile time size implement the `Sized` trait, raw slices do **not** implement it
- Slices must always be behind a reference type, i.e. `&[T]` and `&mut [T]` (but also `Box<[T]>` etc)
- The length of the slice is always stored together with the reference



# Creating slices

Because we cannot create slices out of thin air, they have to be located somewhere. There are three possible ways to create slices:

- Using a borrow
  - We can borrow from arrays and vectors to create a slice of their entire contents
- Using ranges
  - We can use ranges to create a slice from parts of a vector or array
- Using a literal (for immutable slices only)
  - We can have memory statically available from our compiled binary

# Creating slices

Using a borrow

```
1  fn sum(data: &[i32]) -> i32 { /* ... */ }
2
3  fn main() {
4      let v = vec![1, 2, 3, 4, 5, 6];
5      let total = sum(&v);
6      println!("{}", total);
7  }
```

# Creating slices

## Using ranges

```
1  fn sum(data: &[i32]) -> i32 { /* ... */ }
2
3  fn main() {
4      let v = vec![0, 1, 2, 3, 4, 5, 6];
5      let all = sum(&v[..]);
6      let except_first = sum(&v[1..]);
7      let except_last = sum(&v[..5]);
8      let except_ends = sum(&v[1..5]);
9  }
```

- The range ``start..end`` contains all values ``x`` with ``start <= x < end``.
- Note: you can also use ranges on their own, for example in a for loop:

```
1  fn main() {
2      for i in 0..10 {
3          println!("{}", i);
4      }
5  }
```

# Creating slices

From a literal

```
1  fn sum(data: &[i32]) -> i32 { todo!("Sum all items in `data`") }
2
3  fn get_v_arr() -> &'static [i32] {
4      &[0, 1, 2, 3, 4, 5, 6]
5  }
6
7  fn main() {
8      let all = sum(get_v_arr());
9  }
```

- Interestingly ``get_v_arr`` works, even though the literal looks like it would only exist temporarily
- Literals actually exist during the entire lifetime of the program
- ``&'static`` here is used to indicate that this slice will exist the entire lifetime of the program

# Strings

We have already seen the `String` type being used before, but let's dive a little deeper

- Strings are used to represent text
- In Rust they are always valid UTF-8
- Their data is stored on the heap
- A `String` is almost the same as `Vec<u8>` with extra checks to prevent creating invalid text



# Strings

Let's take a look at some strings

```
1  fn main() {
2      let s = String::from("Hello world\nSee you!");
3      println!("{:?}", s.split_once(" "));
4      println!("{}", s.len());
5      println!("{:?}", s.starts_with("Hello"));
6      println!("{}", s.to_uppercase());
7      for line in s.lines() {
8          println!("{}", line);
9      }
10 }
```

# String literals

We have already seen string literals being used while constructing a string. The string literal is what arrays are to vectors

```
1  fn main() {  
2      let s1 = "Hello world";  
3      let s2 = String::from("Hello world");  
4  }
```

# String literals

We have already seen string literals being used while constructing a string. The string literal is what arrays are to vectors

```
1  fn main() {  
2      let s1: &'static str = "Hello world";  
3      let s2: String = String::from("Hello world");  
4  }
```

- `s1` is actually a slice, a string slice

# String literals

We have already seen string literals being used while constructing a string. The string literal is what arrays are to vectors

```
1  fn main() {  
2      let s1: &str = "Hello world";  
3      let s2: String = String::from("Hello world");  
4  }
```

- `s1` is actually a slice, a string slice

## ``str`` - the string slice

It should be possible to have a reference to part of a string. But what is it?

- Not ``[u8]`` : not every sequence of bytes is valid UTF-8
- Not ``[char]`` : we could not create a slice from a string since it is stored as UTF-8 encoded bytes
- We introduce a new special kind of slice: ``str``
- For string slices we do not use brackets!

``str``, ``String``, ``[T; N]``, ``Vec``

Static

Dynamic

Borrowed

``[T; N]``

``Vec<T>``

``&[T]``

-

``String``

``&str``

- There is no static variant of `str`
- This would only be useful if we wanted strings of an exact length
- But just like we had the static slice literals, we can use ``&'static str`` literals for that instead!

# `String` or `str`

When do we use `String` and when do we use `str`?

```
1 fn string_len(data: &String) -> usize {  
2     data.len()  
3 }
```

# `String` or `str`

When do we use `String` and when do we use `str`?

```
1 fn string_len(data: &str) -> usize {  
2     data.len()  
3 }
```

- Prefer `&str` over `String` whenever possible
- If you need to mutate a string you might try `&mut str`, but you cannot change a slice's length
- Use `String` or `&mut String` if you need to fully mutate the string




Smart pointers

# Put it in a `Box`

That pointer from the stack to the heap, how do we create such a thing?

- Boxing something is the way to store a value on the heap
- A `Box` uniquely owns that value, there is no one else that also owns that same value
- Even if the type inside the box is `Copy`, the box itself is not, move semantics apply to a box.

```
1 fn main() {  
2     // put an integer on the heap  
3     let boxed_int = Box::new(10);  
4 }
```



# Boxing

There are several reasons to box a variable on the heap

- When something is too large to move around
- We need something that is sized dynamically
- For writing recursive data structures

```
1  struct Node {  
2      data: Vec<u8>,  
3      parent: Box<Node>,  
4  }
```



# Summary